

✓ PES University, Bangalore

Established under Karnataka Act No. 16 of 2013

✓ UE22AM343AB4 - Advanced Data Analytics

Designed by Sathwik HJ

Student Details

- Name : **Shusrith S**
- SRN : **PES1UG22AM155**

✓ Data cleaning and Preprocessing

✓ Context

As an analyst at Torque Titans, you've been given an exciting opportunity to work with a comprehensive dataset that spans the motorcycle market from 1894 to 2022. Your primary responsibility is to clean and preprocess this data to ensure its quality and readiness for analysis. By doing so, you'll enable your team to extract valuable insights that will drive Torque Titans forward in a competitive market. This critical task will set the foundation for innovative, data-driven strategies that will fuel the company's success in the industry.

Let's dive in!

```
!wget https://raw.githubusercontent.com/MBUYt0n/ada/refs/heads/main/ADA_Workshe
```

```
➡ --2024-09-20 17:56:38-- https://raw.githubusercontent.com/MBUYt0n/ada/refs
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199
HTTP request sent, awaiting response... 200 OK
Length: 32644306 (31M) [text/plain]
Saving to: 'all_bikes.csv.1'
```

```
all_bikes.csv.1      100%[=====>]  31.13M   183MB/s   in 0.2s
```

```
2024-09-20 17:56:38 (183 MB/s) - 'all_bikes.csv.1' saved [32644306/32644306]
```

... About the dataset

▼ ABOUT the dataset

- "all_bikes.csv"
- Each record of the dataset represents a bike model which contains whereas details about it.

```
%pip install matplotlib pandas
```

```
%pip install numpy
```

```
%pip install scikit-learn
```



```
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.1
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/di
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.10/dis
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.1
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/pytho
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/di
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/di
Requirement already satisfied: numpy<2.0,>=1.17.3 in /usr/local/lib/python3
Requirement already satisfied: scipy>=1.5.0 in /usr/local/lib/python3.10/di
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/d
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/pytho
```

```
# Step 1: Import the required modules
```

```
import pandas as pd
```

```
import numpy as np
```

```
# Load the dataset
```

```
df = pd.read_csv('all_bikes.csv') # Replace with your dataset path
```



```
<ipython-input-31-ef01ff8d2199>:6: DtypeWarning: Columns (80,81) have mixed
df = pd.read_csv('all_bikes.csv') # Replace with your dataset path
```

▼ Note:

Give reasons/explanations/reasoning for each question

Step 1: Understanding the Dataset Start by closely examining the dataset to understand its **structure**. Are there any discrepancies or inconsistencies? Are there columns that may not provide valuable insights? Begin by importing the required modules, such as Pandas, and let's gather some initial insights from the data:

- Analyze the number of columns, data types, and the number of values in each column.
- Calculate basic statistics like averages, minimums, and maximums for numerical data.

```
#Pandas Describe
stats = df.describe(include='all')
stats
```

	Model	Year	Category	Rating	Displacement	Engine type	Engin detail
count	38472	38472.000000	38472	38472	37461	38461	639
unique	18597	NaN	18	255	1330	30	130
top	Harley-Davidson Servi-Car GE	NaN	Scooter	Do you know this bike? Click here to rate it. W...	125.0 ccm (7.63 cubic inches)	Single cylinder, four-stroke	Titanium valve
freq	38	NaN	6669	13018	1481	14703	16
mean	NaN	2003.195883	NaN	NaN	NaN	NaN	NaN
std	NaN	20.083372	NaN	NaN	NaN	NaN	NaN
min	NaN	1894.000000	NaN	NaN	NaN	NaN	NaN
25%	NaN	2000.000000	NaN	NaN	NaN	NaN	NaN
50%	NaN	2010.000000	NaN	NaN	NaN	NaN	NaN
75%	NaN	2016.000000	NaN	NaN	NaN	NaN	NaN
max	NaN	2022.000000	NaN	NaN	NaN	NaN	NaN

11 rows × 85 columns

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 38472 entries, 0 to 38471
Data columns (total 85 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Model                                38472 non-null  object
1   Year                                38472 non-null  int64
2   Category                            38472 non-null  object
3   Rating                              38472 non-null  object
4   Displacement                        37461 non-null  object
5   Engine type                         38461 non-null  object
6   Engine details                      6390 non-null   object
_   _
```

7	Power	26110	non-null	object
8	Top speed	12730	non-null	object
9	Bore x stroke	28689	non-null	object
10	Fuel system	27844	non-null	object
11	Ignition	16529	non-null	object
12	Cooling system	34258	non-null	object
13	Gearbox	32675	non-null	object
14	Transmission type	32861	non-null	object
15	Driveline	8324	non-null	object
16	Frame type	17303	non-null	object
17	Front suspension	26107	non-null	object
18	Rear suspension	25540	non-null	object
19	Wheels	9387	non-null	object
20	Seat	4846	non-null	object
21	Dry weight	22483	non-null	object
22	Power/weight ratio	15213	non-null	object
23	Clutch	15152	non-null	object
24	Overall width	18738	non-null	object
25	Fuel capacity	31704	non-null	object
26	Comments	13700	non-null	object
27	Exhaust system	5900	non-null	object
28	Compression	23405	non-null	object
29	Fuel control	22008	non-null	object
30	Lubrication system	10158	non-null	object
31	Front tire	31982	non-null	object
32	Rear tire	32008	non-null	object
33	Front brakes	36889	non-null	object
34	Rear brakes	36696	non-null	object
35	Weight incl. oil, gas, etc	14994	non-null	object
36	Overall length	22242	non-null	object
37	Ground clearance	14205	non-null	object
38	Wheelbase	25493	non-null	object
39	Oil capacity	3786	non-null	object
40	Color options	24327	non-null	object
41	Starter	26845	non-null	object
42	Electrical	3085	non-null	object
43	Valves per cylinder	16173	non-null	float64
44	Diameter	18816	non-null	object
45	Carrying capacity	2702	non-null	object
46	Modifications compared to previous model	344	non-null	object
47	Seat height	24182	non-null	object
48	Overall height	16635	non-null	object
49	Fuel consumption	6176	non-null	object
50	Greenhouse gases	6176	non-null	object
51	Torque	16634	non-null	object
52	Max RPM	663	non-null	float64

Step 2: Handling Missing Values Next, let's address missing values in the dataset. For numeric columns, we'll fill the missing values with the column's average. For categorical or other columns, choose an appropriate strategy based on what makes sense for the data—this could be filled with a placeholder like "Unknown" or the most frequent value.

```
null_percentage = df.isna().sum() / len(df)
null_percentage
```

Model	0.000000
Year	0.000000
Category	0.000000
Rating	0.000000
Displacement	0.026279
...	...
Belt width	0.999064
Pulley teeth	0.999116
Chain size	0.996985
Factory warranty	0.859196
Service interval	0.998414

85 rows × 1 columns

dtype: float64

```
a = df.columns[np.where(df.isna().sum() > 0)]
a
```

```
Index(['Displacement', 'Engine type', 'Engine details', 'Power', 'Top
speed',
      'Bore x stroke', 'Fuel system', 'Ignition', 'Cooling system',
      'Gearbox',
      'Transmission type', 'Driveline', 'Frame type', 'Front suspension',
      'Rear suspension', 'Wheels', 'Seat', 'Dry weight', 'Power/weight
ratio',
      'Clutch', 'Overall width', 'Fuel capacity', 'Comments',
      'Exhaust system', 'Compression', 'Fuel control', 'Lubrication
system',
      'Front tire', 'Rear tire', 'Front brakes', 'Rear brakes',
      'Weight incl. oil, gas, etc', 'Overall length', 'Ground clearance',
      'Wheelbase', 'Oil capacity', 'Color options', 'Starter',
      'Electrical',
      'Valves per cylinder', 'Diameter', 'Carrying capacity',
      'Modifications compared to previous model', 'Seat height',
      'Overall height', 'Fuel consumption', 'Greenhouse gases', 'Torque',
      'Max RPM', 'Light', 'Alternate seat height', 'Rake (fork angle)',
      '0-100 km/h (0-62 mph)', 'Front wheel travel', 'Rear wheel travel',
      'Engine oil', 'Instruments', '60-140 km/h (37-87 mph), highest
gear',
      'Front percentage of weight', 'Trail', 'Brake fluid', 'Coolant',
      'Spark plugs', 'Idle speed', 'Tire pressure front',
      'Tire pressure rear', 'Fork tube size', 'Chain links', 'Sprockets',
      'Reserve fuel capacity', '1/4 mile (0.4 km)', 'Emission details',
      'Rear percentage of weight', 'Oil filter', 'Battery', 'Belt teeth',
      'Belt width', 'Pulley teeth', 'Chain size', 'Factory warranty',
      'Service interval'],
      dtype='object')
```

```
dtype= object ,
```

```
numerical = df.describe()
means = numerical.loc["mean"]
df[numerical.columns] = df[numerical.columns].fillna(means)
df[numerical.columns].isna().sum()
```

	0
Year	0
Valves per cylinder	0
Max RPM	0
Front percentage of weight	0
Chain links	0
Rear percentage of weight	0
Belt teeth	0
Chain size	0

dtype: int64

```
non_numeric = list(set(a) - set(numerical.columns))
non_numeric_df = df[non_numeric]
non_numeric_modes = non_numeric_df.describe().loc["top"]
non_numeric_modes
```

	top
Seat	Dual seat
Fuel control	Double Overhead Cams/Twin Cam (DOHC)
Tire pressure rear	36 PSI (2.5 Bar or 250 kPa)
Engine type	Single cylinder, four-stroke
Reserve fuel capacity	4.00 litres (1.06 US gallons)
...	...
Top speed	45.0 km/h (28.0 mph)
Greenhouse gases	129.9 CO2 g/km. (CO2 - Carbon dioxide emission)
Spark plugs	NGK DCPR7E, NGK DCPR7EIX
Driveline	CVT
Lubrication system	Wet sump

74 rows × 1 columns

.. .. .

dtype: object

```
df[non_numeric] = df[non_numeric].fillna(non_numeric_modes)
df.isna().sum()
```

	0
Model	0
Year	0
Category	0
Rating	0
Displacement	0
...	...
Belt width	0
Pulley teeth	0
Chain size	0
Factory warranty	0
Service interval	0

85 rows × 1 columns

dtype: int64

Step 3: Eliminating Redundancies Take a look at the 0-100 column—do we really need speed in two different units? Let's clean up this redundancy. Be mindful, though; this might not be the only column with unnecessary duplication.

```
df.columns
```

```
Index(['Model', 'Year', 'Category', 'Rating', 'Displacement', 'Engine
type',
      'Engine details', 'Power', 'Top speed', 'Bore x stroke', 'Fuel
system',
      'Ignition', 'Cooling system', 'Gearbox', 'Transmission type',
      'Driveline', 'Frame type', 'Front suspension', 'Rear suspension',
      'Wheels', 'Seat', 'Dry weight', 'Power/weight ratio', 'Clutch',
      'Overall width', 'Fuel capacity', 'Comments', 'Exhaust system',
      'Compression', 'Fuel control', 'Lubrication system', 'Front tire',
      'Rear tire', 'Front brakes', 'Rear brakes',
      'Weight incl. oil, gas, etc', 'Overall length', 'Ground clearance',
      'Wheelbase', 'Oil capacity', 'Color options', 'Starter',
      'Electrical',
      'Valves per cylinder', 'Diameter', 'Carrying capacity',
      'Modifications compared to previous model', 'Seat height',
      'Overall height', 'Fuel consumption', 'Greenhouse gases', 'Torque']
)
```

```

        'Overall height', 'Fuel consumption', 'Greenhouse gases', 'Torque',
        'Max RPM', 'Light', 'Alternate seat height', 'Rake (fork angle)',
        '0-100 km/h (0-62 mph)', 'Front wheel travel', 'Rear wheel travel',
        'Engine oil', 'Instruments', '60-140 km/h (37-87 mph), highest
gear',
        'Front percentage of weight', 'Trail', 'Brake fluid', 'Coolant',
        'Spark plugs', 'Idle speed', 'Tire pressure front',
        'Tire pressure rear', 'Fork tube size', 'Chain links', 'Sprockets',
        'Reserve fuel capacity', '1/4 mile (0.4 km)', 'Emission details',
        'Rear percentage of weight', 'Oil filter', 'Battery', 'Belt teeth',
        'Belt width', 'Pulley teeth', 'Chain size', 'Factory warranty',
        'Service interval'],
        dtype='object')

```

```

a = len(df) - len(df[df["Tire pressure rear"] == df["Tire pressure front"]])
b = len(df) - len(df[df["Front brakes"] == df["Rear brakes"]])
c = len(df) - len(df[df["Rear percentage of weight"] == df["Front percentage of
d = len(df) - len(df[df["Front wheel travel"] == df["Rear wheel travel"]])
print(a, b, c, d)

```

156 21499 38467 36349

```

a = df[["Gearbox", "Transmission type"]]
a.value_counts()

```

		count
Gearbox	Transmission type	
6-speed	Chain (final drive)	14595
5-speed	Chain (final drive)	7915
Automatic	Belt (final drive)	3828
4-speed	Chain (final drive)	2258
5-speed	Shaft drive (cardan) (final drive)	1990
6-speed	Belt (final drive)	1974
Automatic	Chain (final drive)	1599
6-speed	Shaft drive (cardan) (final drive)	1132
5-speed	Belt (final drive)	1087
Automatic	Shaft drive (cardan) (final drive)	595
4-speed	Shaft drive (cardan) (final drive)	419
3-speed	Chain (final drive)	406
1-speed	Chain (final drive)	224
4-speed	Belt (final drive)	119
3-speed	Shaft drive (cardan) (final drive)	57
2-speed	Chain (final drive)	53

1-speed	Belt (final drive)	51
4-speed with reverse	Shaft drive (cardan) (final drive)	49
2-speed	Shaft drive (cardan) (final drive)	25
	Belt (final drive)	15
3-speed	Belt (final drive)	15
7-speed	Shaft drive (cardan) (final drive)	13
	Chain (final drive)	11
100-speed	Belt (final drive)	7
	Shaft drive (cardan) (final drive)	6
2-speed automatic	Shaft drive (cardan) (final drive)	5
1-speed	Shaft drive (cardan) (final drive)	5
5-speed with reverse	Shaft drive (cardan) (final drive)	5
2-speed automatic	Chain (final drive)	3
10-speed	Shaft drive (cardan) (final drive)	3
8-speed	Chain (final drive)	3
100-speed	Chain (final drive)	2
3-speed automatic	Chain (final drive)	1
6-speed with reverse	Shaft drive (cardan) (final drive)	1
10-speed	Chain (final drive)	1

dtype: int64

```
df.drop(
    [
        "Weight incl. oil, gas, etc",
        "1/4 mile (0.4 km)",
        "60-140 km/h (37-87 mph), highest gear",
        "Tire pressure front",
        "Greenhouse gases"
    ], axis=1, inplace=True
)
```

Step 4: Duplicates Now, check for any duplicate records in the dataset. If duplicates are found, remove them to avoid any skewed analysis.

```
df.drop_duplicates(inplace=True)
```

Step 5: Engine Details Column Examine the **Engine Details** column carefully. Will this column be useful in providing insights, or is it redundant or irrelevant to the analysis? Decide whether to keep or drop it.

```
df["Engine details"].value_counts() / len(df)
```

	count
Engine details	
Titanium valves	0.838090
Reed intake.	0.003613
90° V-twin	0.002495
Reed valve.	0.002079
Balancer shaft	0.002027
...	...
Permanent magnet synchronous motor in a disc armature design.	0.000026
Permanent magnet synchronous motorin a disc armature design	0.000026
48 V BLDC motor with outer rotor	0.000026
16 valves with variable valve timing	0.000026
Fuel injection: ø42 mm x 2	0.000026

1301 rows × 1 columns

dtype: float64

```
df.drop("Engine details", axis=1, inplace=True)
```

Start coding or [generate](#) with AI.

Step 6: Preparing for Future Text Processing Torque Titans might explore text processing on some of the data in the future, so let's be proactive! We can tokenize the strings in the relevant columns to ensure we're ready for text analysis down the line. This involves splitting text into individual tokens (words) and storing them for future use.

There are a lot of tokeniser available, note: Torque Titans are potentially looking to integrate with openai.

```
!pip install tiktoken
```

```
Requirement already satisfied: tiktoken in /usr/local/lib/python3.10/dist-p
Requirement alreadyv satisfied: regex>=2022.1.18 in /usr/local/lib/pvthon3.1
```

Requirement already satisfied: requests>=2.26.0 in /usr/local/lib/python3.1
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/p
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/di
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3

```
import tiktoken
```

```
enc = tiktoken.get_encoding("o200k_base")
```

```
non_numeric = df.select_dtypes(exclude=np.number)
```

```
tokenized_df = non_numeric.apply(lambda x: x.apply(lambda y: enc.encode(y)))
```

```
tokenized_df
```

	Model	Category	Rating	Displacement	Engine type	Power	Top speed	Bore x stroke
0	[32,	[4764, 2884, 820, 1277, 12086]	[220,	[3796, 13, 24, 274, 7871, 350, 18, 13, 3000, 6...	[13301, 49938, 11, 1920, 6321, 11472]	[24, 13, 20, 21979, 350, 21, 13, 24, 2029, 220, 372,...	[2548, 13, 15, 8571, 14174, 350, 2029, 13, 15,...	[1723, 13, 15, 1215, 220, 3255, 13, 22, 8957, ...
	41166,		18, 13,					
	128298,		22,					
	47, 220,		220,					
1	1434,	[50837]	5310,	[3796, 13, 24, 274, 7871, 350, 18, 13, 3000, 6...	[13301, 49938, 11, 1920, 6321, 11472]	[24, 13, 20, 21979, 350, 21, 13, 24, 2029, 220, 372,...	[2548, 13, 15, 8571, 14174, 350, 2029, 13, 15,...	[1723, 13, 15, 1215, 220, 3255, 13, 22, 8957, ...
	6437,		4383,					
	2884]		842,					
			290,					
2		[17260, 3598, 597]	1...	[3796, 13, 24, 274, 7871, 350, 18, 13, 3000, 6...	[13301, 49938, 11, 1920, 6321, 11472]	[24, 13, 20, 21979, 350, 21, 13, 24, 2029, 220, 372,...	[2548, 13, 15, 8571, 14174, 350, 2029, 13, 15,...	[1723, 13, 15, 1215, 220, 3255, 13, 22, 8957, ...
	[32,		[220,					
	41166,		18, 13,					
	128298,		21,					
3	47, 220,	[4764, 2884, 820, 1277, 12086]	220,	[16059, 13, 15, 274, 7871, 350, 22, 13, 5085, ...	[13301, 49938, 11, 4242, 6321, 11472]	[899, 13, 20, 21979, 350, 24, 13, 16, 220, 372...	[2548, 13, 15, 8571, 14174, 350, 2029, 13, 15,...	[6733, 13, 15, 1215, 220, 6733, 13, 15, 8957, ...
	1434,		5310,					
	8220,		4383,					
	46121...		842,					

			[6449, 481, 1761, 495, 17431, 30, 3524, 2105, ...			[13301, 49938, 11, 4242, 6321, 11472]	[899, 13, 20, 21979, 350, 24, 13, 16, 220, 372...	[2548, 13, 15, 8571, 14174, 350, 2029, 13, 15,...	[6733, 13, 15, 1215, 220, 6733, 13, 15, 8957, ...
4	[32, 41166, 11448, 18, 220, 10676, 6437, 2884,...	[4764, 2884, 820, 1277, 12086]			[16059, 13, 15, 274, 7871, 350, 22, 13, 5085, ...				
...
38467	[347, 1191, 15928, 5239, 394, 336, 220, 4095, ...	[188210, 8708]	[6449, 481, 1761, 495, 17431, 30, 3524, 2105, ...		[10676, 13, 15, 274, 7871, 350, 22, 13, 8876, ...	[78080]	[19, 13, 15, 21979, 350, 17, 13, 24, 220, 372,...	[3898, 13, 15, 8571, 14174, 350, 1723, 13, 19,...	[6733, 13, 15, 1215, 220, 6733, 13, 15, 8957, ...
38468	[347, 1191, 15928, 5239, 394, 336, 220, 3234, 15]	[188210, 8708]	[6449, 481, 1761, 495, 17431, 30, 3524, 2105, ...		[10676, 13, 15, 274, 7871, 350, 22, 13, 8876, ...	[78080]	[21, 13, 22, 21979, 350, 19, 13, 24, 220, 372,...	[10116, 13, 15, 8571, 14174, 350, 6231, 13, 15...	[6733, 13, 15, 1215, 220, 6733, 13, 15, 8957, ...
38469	[823, 130874, 113621]	[188210, 8708]	[6449, 481, 1761, 495, 17431, 30, 3524, 2105, ...		[10676, 13, 15, 274, 7871, 350, 22, 13, 8876, ...	[78080]	[17, 13, 22, 21979, 350, 17, 13, 15, 220, 372,...	[2548, 13, 15, 8571, 14174, 350, 2029, 13, 15,...	[6733, 13, 15, 1215, 220, 6733, 13, 15, 8957, ...
38470	[823, 130874, 122050]	[188210, 8708]	[6449, 481, 1761, 495, 17431, 30, 3524, 2105, ...		[10676, 13, 15, 274, 7871, 350, 22, 13, 8876, ...	[78080]	[17, 13, 22, 21979, 350, 17, 13, 15, 220, 372,...	[2548, 13, 15, 8571, 14174, 350, 2029, 13, 15,...	[6733, 13, 15, 1215, 220, 6733, 13, 15, 8957, ...
38471	[823, 130874, 45558]	[188210, 8708]	[6449, 481, 1761, 495, 17431, 30, 3524, 2105,		[10676, 13, 15, 274, 7871, 350, 22, 13, 8876, ...	[78080]	[17, 13, 22, 21979, 350, 17, 13, 15, 220, 372	[2548, 13, 15, 8571, 14174, 350, 2029, 13, 15	[6733, 13, 15, 1215, 220, 6733, 13, 15, 8957,

38472 rows × 71 columns

df.columns

```
Index(['Model', 'Year', 'Category', 'Rating', 'Displacement', 'Engine
type',
      'Power', 'Top speed', 'Bore x stroke', 'Fuel system', 'Ignition',
      'Cooling system', 'Gearbox', 'Transmission type', 'Driveline',
      'Frame type', 'Front suspension', 'Rear suspension', 'Wheels',
      'Seat',
      'Dry weight', 'Power/weight ratio', 'Clutch', 'Overall width',
      'Fuel capacity', 'Comments', 'Exhaust system', 'Compression',
      'Fuel control', 'Lubrication system', 'Front tire', 'Rear tire',
      'Front brakes', 'Rear brakes', 'Overall length', 'Ground
clearance',
      'Wheelbase', 'Oil capacity', 'Color options', 'Starter',
      'Electrical',
      'Valves per cylinder', 'Diameter', 'Carrying capacity',
      'Modifications compared to previous model', 'Seat height',
      'Overall height', 'Fuel consumption', 'Torque', 'Max RPM', 'Light',
      'Alternate seat height', 'Rake (fork angle)', '0-100 km/h (0-62
mph)',
      'Front wheel travel', 'Rear wheel travel', 'Engine oil',
      'Instruments',
      'Front percentage of weight', 'Trail', 'Brake fluid', 'Coolant',
      'Spark plugs', 'Idle speed', 'Tire pressure rear', 'Fork tube
size',
      'Chain links', 'Sprockets', 'Reserve fuel capacity', 'Emission
details',
      'Rear percentage of weight', 'Oil filter', 'Battery', 'Belt teeth',
      'Belt width', 'Pulley teeth', 'Chain size', 'Factory warranty',
      'Service interval'],
      dtype='object')
```

Step 7: Designing a Rating System Here comes the fun part—Torque Titans wants to roll out a rating system for their motorcycles! You'll take into account the following factors:

- Speed
- Engine type (feel free to assign weight based on your opinion of which engines are superior)
- 0-100 acceleration
- Power
- Torque
- Weight
- RPM (you might want to cross-reference with the torque column).

Using these factors, create a new column called Rating, which is scaled from 1 to 4 (no

decimals). You don't need to give all columns equal weight—experiment to find the ideal balance for what you think makes a great motorcycle!

Make sure to normalise the values before considering them for the rating column, as they can add a bias to the calculations, Lets try using MaxAbsScaler and try keep the values between [-1,1]

```
A = ['Engine type', 'Top speed', 'Power', 'Torque', 'Dry weight']
features = df[A]
```

```
for i in A:
    print(features[i].value_counts())
```

```
Engine type
Single cylinder, four-stroke    14714
V2, four-stroke                 7405
Single cylinder, two-stroke     5982
In-line four, four-stroke       3152
Twin, four-stroke               2888
Electric                       979
Two cylinder boxer, four-stroke 862
In-line three, four-stroke      794
Twin, two-stroke                500
V4, four-stroke                 452
Six cylinder boxer, four-stroke 137
In-line six, four-stroke        111
V8, four-stroke                 79
Two cylinder boxer, two-stroke  74
In-line three, two-stroke       58
Four cylinder boxer, four-stroke 52
V6, four-stroke                 42
V2, two-stroke                  39
Diesel                          37
Square four cylinder            33
Gas turbine                     19
In-line four, two-stroke        17
Dual disk Wankel                13
Radial                          10
Single disk Wankel               8
V4, two-stroke                  7
In-line six, two-stroke         3
V3, two-stroke                  3
V10, four-stroke                1
Four cylinder boxer, two-stroke 1
Name: count, dtype: int64

Top speed
45.0 km/h (28.0 mph)    26278
90.0 km/h (55.9 mph)    445
100.0 km/h (62.1 mph)   422
110.0 km/h (68.4 mph)   401
95.0 km/h (59.0 mph)    342
...
267.0 km/h (165.9 mph)   1
67.0 km/h (41.6 mph)     1
224.5 km/h (139.5 mph)   1
322 km/h (200.0 mph)     1
```

```

32.2 km/h (20.0 mph)      1
133.6 km/h (83.0 mph)     1
Name: count, Length: 416, dtype: int64
Power
27.0 HP (19.7 kW)) @ 6000 RPM    12458
50.0 HP (36.5 kW)) @ 6500 RPM      90
27.0 HP (19.7 kW)) @ 6500 RPM      85
15.0 HP (10.9 kW))                83
2.7 HP (2.0 kW))                  76
...
65.7 HP (48.0 kW)) @ 4700 RPM      1
71.0 HP (51.8 kW)) @ 5500 RPM      1
63.0 HP (46.0 kW)) @ 4700 RPM      1
64.0 HP (46.7 kW)) @ 4900 RPM      1
136.0 HP (99.3 kW))                1
Name: count, Length: 4914, dtype: int64

```

```

# Remove commas before converting to float
features["Top speed"] = features["Top speed"].apply(lambda x: float(x.split()[0].replace(',', '')))
features["Dry weight"] = features["Dry weight"].apply(lambda x: float(x.split()[0].replace(',', '')))
features["RPM"] = features["Torque"].apply(lambda x: float(x.split()[-2].replace(',', '')))
features["Power"] = features["Power"].apply(lambda x: float(x.split()[0].replace(',', '')))
features["Torque"] = features["Torque"].apply(lambda x: float(x.split()[0].replace(',', '')))

```

```

<ipython-input-82-390fbd474cf6>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs
features["Top speed"] = features["Top speed"].apply(lambda x: float(x.split()[0].replace(',', '')))
<ipython-input-82-390fbd474cf6>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs
features["Dry weight"] = features["Dry weight"].apply(lambda x: float(x.split()[0].replace(',', '')))
<ipython-input-82-390fbd474cf6>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs
features["RPM"] = features["Torque"].apply(lambda x: float(x.split()[-2].replace(',', '')))
<ipython-input-82-390fbd474cf6>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs
features["Power"] = features["Power"].apply(lambda x: float(x.split()[0].replace(',', '')))
<ipython-input-82-390fbd474cf6>:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs
features["Torque"] = features["Torque"].apply(lambda x: float(x.split()[0].replace(',', '')))

```

features

	Engine type	Top speed	Power	Torque	Dry weight	RPM			
0	Single cylinder, two-stroke	45.0	9.5	52.0	78.0	4000.0			
1	Single cylinder, two-stroke	45.0	9.5	52.0	78.0	4000.0			
2	Single cylinder, two-stroke	45.0	9.5	52.0	78.0	4000.0			
3	Single cylinder, four-stroke	45.0	12.5	8.5	110.0	8000.0			
4	Single cylinder, four-stroke	45.0	12.5	8.5	110.0	8000.0			
...			
38467	Electric	65.0	4.0	52.0	144.0	4000.0			
38468	Electric	82.0	6.7	52.0	130.0	4000.0			
38469	Electric	45.0	2.7	52.0	110.0	4000.0			
38470	Electric	45.0	2.7	52.0	110.0	4000.0			

Next steps:

Generate code with features

☐ View recommended plots

New interactive sheet

```
for i in features:
    print(features[i].value_counts())

Engine type
Single cylinder, four-stroke    14714
V2, four-stroke                 7405
Single cylinder, two-stroke     5982
In-line four, four-stroke       3152
Twin, four-stroke               2888
Electric                       979
Two cylinder boxer, four-stroke 862
In-line three, four-stroke      794
Twin, two-stroke                500
V4, four-stroke                 452
Six cylinder boxer, four-stroke 137
In-line six, four-stroke        111
V8, four-stroke                 79
Two cylinder boxer, two-stroke  74
In-line three, two-stroke       58
Four cylinder boxer, four-stroke 52
V6, four-stroke                 42
V2, two-stroke                  39
Diesel                          37
Square four cylinder            33
Gas turbine                     19
In-line four, two-stroke        17
Dual disk Wankel                13
Radial                          10
Single disk Wankel               8
V4, two-stroke                  7
```



```

In-line six, two-stroke      3
V3, two-stroke               3
V10, four-stroke             1
Four cylinder boxer, two-stroke 1
Name: count, dtype: int64
Top speed
45.0      26278
90.0       445
100.0      422
110.0      401
95.0       342

...
267.0      1
67.0       1
224.5      1
32.2       1
133.6      1
Name: count, Length: 416, dtype: int64
Power
27.0      13207
50.0       582
17.0       512
100.0      455
15.0       369

...
40.4      1
82.9      1
93.8      1
19.1      1
194.5     1
Name: count, Length: 849, dtype: int64

```




```

d = {"Single" : 1, "single": 1, "two" : 2, "three" : 3, "four" : 4, "five" : 5,
e = features["Engine type"]
s = []
for i in e:
    c = 0
    p = i.split(",")
    for j in d:
        for k in p:
            if j in k:
                c += d[j]
    s.append(c)
features["Engine type"] = s
features

```

<ipython-input-85-125ba2a9ffae>:12: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs>
features["Engine type"] = s

	Engine type	Top speed	Power	Torque	Dry weight	RPM	
0	3	45.0	9.5	52.0	78.0	4000.0	
1	3	45.0	9.5	52.0	78.0	4000.0	

2	3	45.0	9.5	52.0	78.0	4000.0
3	5	45.0	12.5	8.5	110.0	8000.0
4	5	45.0	12.5	8.5	110.0	8000.0
...
38467	0	65.0	4.0	52.0	144.0	4000.0
38468	0	82.0	6.7	52.0	130.0	4000.0
38469	0	45.0	2.7	52.0	110.0	4000.0
38470	0	45.0	2.7	52.0	110.0	4000.0
38471	0	45.0	2.7	52.0	110.0	4000.0

38472 rows × 6 columns

Next steps:

Generate code with features




☒
View recommended plots

New interactive sheet

```

from sklearn.preprocessing import MaxAbsScaler
scaler = MaxAbsScaler()
normalized_df = pd.DataFrame(scaler.fit_transform(features), columns=features.co
normalized_df

```

	Engine type	Top speed	Power	Torque	Dry weight	RPM	
0	0.3	0.069231	0.011807	0.073034	0.078	0.275862	
1	0.3	0.069231	0.011807	0.073034	0.078	0.275862	
2	0.3	0.069231	0.011807	0.073034	0.078	0.275862	
3	0.5	0.069231	0.015536	0.011938	0.110	0.551724	
4	0.5	0.069231	0.015536	0.011938	0.110	0.551724	
...	
38467	0.0	0.100000	0.004971	0.073034	0.144	0.275862	
38468	0.0	0.126154	0.008327	0.073034	0.130	0.275862	
38469	0.0	0.069231	0.003356	0.073034	0.110	0.275862	
38470	0.0	0.069231	0.003356	0.073034	0.110	0.275862	
38471	0.0	0.069231	0.003356	0.073034	0.110	0.275862	

38472 rows × 6 columns

Next steps:

Generate code with normalized_df

☒
View recommended plots

New interactive sheet

```
weights = [0.8, 0.8, 0.8, 0.8, -0.8, 0.8]
def map_number(n):
    if n < 0.25:
        return 1
    elif n < 0.5:
        return 2
    elif n < 0.75:
        return 3
    else:
        return 4

normalized_df["Rating"] = (normalized_df * weights).sum(axis=1)
normalized_df["Rating"] = normalized_df["Rating"].apply(map_number)
normalized_df
```

	Engine type	Top speed	Power	Torque	Dry weight	RPM	Rating
0	0.3	0.069231	0.011807	0.073034	0.078	0.275862	3
1	0.3	0.069231	0.011807	0.073034	0.078	0.275862	3
2	0.3	0.069231	0.011807	0.073034	0.078	0.275862	3
3	0.5	0.069231	0.015536	0.011938	0.110	0.551724	4
4	0.5	0.069231	0.015536	0.011938	0.110	0.551724	4
...
38467	0.0	0.100000	0.004971	0.073034	0.144	0.275862	1
38468	0.0	0.126154	0.008327	0.073034	0.130	0.275862	2
38469	0.0	0.069231	0.003356	0.073034	0.110	0.275862	1
38470	0.0	0.069231	0.003356	0.073034	0.110	0.275862	1
38471	0.0	0.069231	0.003356	0.073034	0.110	0.275862	1

38472 rows × 7 columns

Next steps:

Generate code with `normalized_df`

 View recommended plots

New interactive sheet

Double-click (or enter) to edit

By completing these steps, you'll not only ensure that the dataset is ready for insightful analysis but also set the stage for exciting innovations at Torque Titans. Let's get started and have fun while we shape the future of motorcycles!

