

✓ PES University, Bangalore

Established under Karnataka Act No. 16 of 2013

UE22AM343AB4 - Advanced Data Analytics

Designed by Nischal H S

Student Details

- Name : **Student Name**
- SRN : **SRN**

✓ ADA Worksheet Part B

Adult Census Income Cleaning and Analysis

Introduction

As a data scientist intern at the U.S. Census Bureau, you've been assigned to clean and analyze the Adult Census Income dataset. Your task is to prepare the data for a machine learning model that will predict whether an individual's annual income exceeds \$50,000. This analysis will inform government policies on education, employment, and economic development. The dataset contains various demographic and socioeconomic factors, but it requires careful preprocessing to ensure accurate results. Your work will involve handling missing data, encoding categorical variables, and performing exploratory data analysis.

First, let's import some of the necessary libraries and load the data.

```
# might make it easier to install the packages directly to ipynb kernel for thi
%pip install pandas numpy matplotlib seaborn scikit-learn scipy imbalanced-lear
```

```
➡ Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-pa
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/di
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: imbalanced-learn in /usr/local/lib/python3.1
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/pyt
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/di
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.1
```

```
Requirement already satisfied: cyclcr>=0.10 in /usr/local/lib/python3.10/d1
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.1
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/d
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/pytho
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-p
```

```
# Note: This assignment might need you to look up syntax, parameters, functions
# pandas: https://pandas.pydata.org/docs/user\_guide/index.html
# numpy: https://numpy.org/doc/stable/user/index.html
# matplotlib: https://matplotlib.org/stable/contents.html
# seaborn: https://seaborn.pydata.org/tutorial.html
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Load the dataset
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.dat"
column_names = ["age", "workclass", "fnlwgt", "education", "education-num", "ma

# TODO: Load the data into a pandas DataFrame named 'df'
# Hint: Use pd.read_csv() with the url and column_names
```

```
df = pd.read_csv(url, names=column_names, sep=",\s", na_values=["?"], engine="p
# Display the first few rows and basic information about the dataset
```

```
df.head()
```



```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   32561 non-null  int64
 1   workclass              30725 non-null  object
 2   fnlwgt                 32561 non-null  int64
 3   education              32561 non-null  object
 4   education-num          32561 non-null  int64
 5   marital-status         32561 non-null  object
 6   occupation              30718 non-null  object
 7   relationship           32561 non-null  object
 8   race                   32561 non-null  object
 9   sex                   32561 non-null  object
10   capital-gain           32561 non-null  int64
11   capital-loss           32561 non-null  int64
12   hours-per-week         32561 non-null  int64
13   native-country         31978 non-null  object
14   income                 32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
df.describe()
```

✓ Question 1: Missing Values

Examine the dataset for missing values. In this dataset, missing values are represented as "?".

a) How many missing values are there in each column?

b) What percentage of the dataset is missing?

```
# TODO: Count the number of missing values in each column
# Hint: Use df.isna() with the right parameter
```

```
# Hint: use df.isin() with the right parameter
df.isna().sum() / len(df) * 100
# TODO: Calculate the percentage of missing values in the entire dataset
```

✓ Question 2: Handling Missing Values

Choose an appropriate strategy to handle missing values in the 'workclass' and 'occupation' columns.

- Explain your chosen strategy and why you think it's appropriate.
- Implement your strategy.

df

```

a = df[['race', 'native-country']]
for i in a.value_counts().items():
    print(i)

(('White', 'United-States'), 25621)
(('Black', 'United-States'), 2832)
(('White', 'Mexico'), 590)
(('Amer-Indian-Eskimo', 'United-States'), 296)
(('Asian-Pac-Islander', 'United-States'), 292)
(('Asian-Pac-Islander', 'Philippines'), 188)
(('Other', 'United-States'), 129)
(('White', 'Germany'), 124)
(('White', 'Canada'), 119)
(('White', 'El-Salvador'), 101)
(('White', 'Cuba'), 90)
(('Asian-Pac-Islander', 'India'), 85)
(('White', 'Puerto-Rico'), 82)
(('White', 'England'), 81)
(('Asian-Pac-Islander', 'South'), 77)
(('Black', 'Jamaica'), 75)
(('Asian-Pac-Islander', 'China'), 73)
(('White', 'Italy'), 73)
(('Asian-Pac-Islander', 'Vietnam'), 65)
(('White', 'Guatemala'), 60)
(('White', 'Poland'), 59)
(('White', 'Columbia'), 51)
(('Asian-Pac-Islander', 'Taiwan'), 48)
(('Black', 'Haiti'), 43)

```

```

(('Black', 'India'), 7)
(('Other', 'Mexico'), 40)
(('White', 'Dominican-Republic'), 39)
(('Asian-Pac-Islander', 'Japan'), 38)
(('White', 'Portugal'), 36)
(('White', 'Iran'), 35)
(('White', 'Peru'), 30)
(('White', 'France'), 28)
(('White', 'Nicaragua'), 28)
(('White', 'Greece'), 28)
(('White', 'Ireland'), 23)
(('Other', 'Puerto-Rico'), 21)
(('White', 'Ecuador'), 19)
(('White', 'Japan'), 19)
(('Asian-Pac-Islander', 'Laos'), 18)
(('Asian-Pac-Islander', 'Cambodia'), 18)
(('Other', 'Dominican-Republic'), 18)
(('Asian-Pac-Islander', 'Hong'), 17)
(('Asian-Pac-Islander', 'Thailand'), 16)
(('Black', 'Trinidad&Tobago'), 16)
(('White', 'Yugoslavia'), 16)
(('White', 'Hungary'), 13)
(('White', 'Scotland'), 12)
(('White', 'Honduras'), 12)
(('Black', 'Dominican-Republic'), 12)
(('Other', 'Ecuador'), 9)
(('Black', 'Puerto-Rico'), 9)
(('Black', 'Germany'), 8)
(('White', 'India'), 8)
(('White', 'Outlying-US(Guam-USVI-etc)'), 8)
(('Black', 'England'), 8)
(('Amer-Indian-Eskimo', 'Mexico'), 8)
(('White', 'Philippines'), 8)
(('Other', 'Columbia'), 7)
(('Asian-Pac-Islander', 'Iran'), 6)

```

```

occupation_dict = {}
a = df[["occupation", "education"]]
for i in list(a.value_counts().items()):
    if i[0][1] not in occupation_dict:
        occupation_dict[i[0][1]] = i[0][0]
occupation_dict

```

```

{'HS-grad': 'Craft-repair',
 'Bachelors': 'Prof-specialty',
 'Some-college': 'Adm-clerical',
 'Masters': 'Prof-specialty',
 'Prof-school': 'Prof-specialty',
 'Doctorate': 'Prof-specialty',
 'Assoc-voc': 'Craft-repair',
 '11th': 'Other-service',
 '10th': 'Other-service',
 'Assoc-acdm': 'Adm-clerical',
 '7th-8th': 'Craft-repair',
 '9th': 'Other-service',
 '12th': 'Other-service',
 '5th-6th': 'Other-service',
 '1st-4th': 'Other-service',
 'Preschool': 'Other-service'}

```

```
    .reset_index(drop=True)
```

```
workclass_dict = {}
a = df[["workclass", "occupation"]]
for i in list(a.value_counts().items()):
    if i[0][1] not in workclass_dict:
        workclass_dict[i[0][1]] = i[0][0]
workclass_dict
```

```
{'Craft-repair': 'Private',
 'Sales': 'Private',
 'Adm-clerical': 'Private',
 'Other-service': 'Private',
 'Exec-managerial': 'Private',
 'Prof-specialty': 'Private',
 'Machine-op-inspct': 'Private',
 'Handlers-cleaners': 'Private',
 'Transport-moving': 'Private',
 'Tech-support': 'Private',
 'Farming-fishing': 'Private',
 'Protective-serv': 'Local-gov',
 'Priv-house-serv': 'Private',
 'Armed-Forces': 'Federal-gov'}
```

```
# TODO: Implement your chosen strategy for handling missing values
# This might involve imputation, removal, or other techniques. Think carefully
```

```
import pandas as pd
```

```
def handle_missing_workclass(df):
    """
    Handle missing values in the 'workclass' column.
    """
    df["workclass"] = df.apply(
        lambda row: (
            "Other"
            if pd.isnull(row["workclass"]) and pd.isnull(row["occupation"])
            else (
                workclass_dict[row["occupation"]]
                if pd.isnull(row["workclass"])
                else row["workclass"]
            )
        ),
        axis=1,
    )
    return df
```

```
def handle_missing_occupation(df):
    """
    Handle missing values in the 'occupation' column.
    """
    df["occupation"] = df.apply(
        lambda row: (
```

```

        lambda row: (
            occupation_dict[row["education"]] if pd.isnull(row["occupation"]) else
        ),
        axis=1
    )
    return df

```

```

def handle_missing_native_country(df):
    """
    Handle missing values in the 'native-country' column.
    """
    df["native-country"].fillna("United-States", inplace=True)
    return df

```

```

# Call the functions to update the DataFrame
df = handle_missing_workclass(df)
df = handle_missing_occupation(df)
df = handle_missing_native_country(df)
df.isna().sum()
# Hint: You could use the mode() method to impute missing values
# Can you go one step ahead and consider the other columns before deciding what

```


✓ Question 3: Categorical Variables

Identify all categorical variables in the dataset.

a) List the categorical variables and their unique values.

b) Are there any categorical variables that have an unusually high number of categories? How might you handle this?

```
# TODO: Identify and list categorical variables
# Hint: Use the select_dtypes() method to identify columns with object dtype
categorical = df.select_dtypes(include=["object"]).columns
categorical
# TODO: Display unique values for each categorical variable
# Hint: You can use the unique() method to get the unique values for each column

Index(['workclass', 'education', 'marital-status', 'occupation',
       'relationship', 'race', 'sex', 'native-country', 'income'],
      dtype='object')

df[categorical].nunique()
```

Native country has very high number of values, we can encode the frequency of the values to show how important they are

```
encoded_df = df[categorical]
```

✓ Question 4: Encoding Categorical Variables

Choose appropriate encoding techniques for the categorical variables. You may use different techniques for different variables based on their characteristics.

- a) Explain your choice of encoding technique for each categorical variable.
- b) Implement the encoding.

```
# TODO: Implement encoding for categorical variables
# This might involve one-hot encoding, label encoding, or other techniques
from sklearn.preprocessing import LabelEncoder, OneHotEncoder

def encode_binary_categorical(df, columns):
    """
    Encode a binary categorical variable using label encoding.
    """
    # Hint: Use the LabelEncoder class from sklearn.preprocessing
    encoder = LabelEncoder()
    for column in columns:
        df[column] = encoder.fit_transform(df[column])
    return df

def encode_multi_categorical(df, columns):
    """
    Encode multi-class categorical variables using one-hot encoding.
    """
    # Hint: Use the OneHotEncoder class from sklearn.preprocessing
    encoder = OneHotEncoder(sparse_output=False, drop='first') # drop='first'
    encoded_columns = encoder.fit_transform(df[columns])
    encoded_df = pd.DataFrame(encoded_columns, columns=encoder.get_feature_name_out())
    df = df.drop(columns, axis=1)
    df = pd.concat([df, encoded_df], axis=1)
    return df

def frequency_encoding(df, columns):
    """
    Encode multi-class categorical
    """
    for column in columns:
        freq = df[column].value_counts(normalize=True)
        df[column] = df[column].map(freq)
    return df

encoded_df = encode_binary_categorical(encoded_df, ["sex", "income"])
encoded_df = encode_multi_categorical(encoded_df, ["race", "workclass", "marital_status"])
encoded_df = frequency_encoding(encoded_df, ["education", "occupation", "native_country"])
encoded_df
```

✓ Question 5: Numerical Variables

Analyze the numerical variables in the dataset.

- a) Create histograms for each numerical variable.
- b) Identify any variables that appear to be skewed. How might you handle this skewness?

```
numerical = df.select_dtypes(include=["int64", "float64"]).columns
numerical

Index(['age', 'fnlwgt', 'education-num', 'capital-gain', 'capital-loss',
      'hours-per-week'],
      dtype='object')

numerical_df = df[numerical]

for i in numerical:
    sns.histplot(data=numerical df. x=i)
```

```
plt.title(i)  
plt.show()
```

```
numerical_df.skew(axis=0, numeric_only=True)
```

```
skewed_vars = abs(numerical_df.skew(axis=0, numeric_only=True)) > 0.5
skewed_vars = skewed_vars[skewed_vars].index
skewed_vars
```

```
Index(['age', 'fnlwtg', 'capital-gain', 'capital-loss'], dtype='object')
```

```
# TODO: Create histograms for numerical variables
```

```
# Hint: Use the histplot() function from the seaborn library
```

```
# TODO: Identify skewed variables and suggest transformations
```

```
# Hint: You can use the skew() method to identify skewed variables
```

```
# For transformations, you could consider using the np.loglp() function
```

```
# Verify the effectiveness of your transformation
```

```
def transform_skewed_variables(df, skewed_vars):
    """
    Apply transformation to skewed numerical variables.
    """
    for var in skewed_vars:
        df[var] = np.loglp(df[var])
    return df
```

```
numerical_df = transform_skewed_variables(numerical_df, skewed_vars)
```

```
<ipython-input-96-1b677f3c0c1f>:16: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs
df[var] = np.loglp(df[var])
```

```
for i in numerical:
    sns.histplot(data=numerical_df, x=i)
    plt.title(i)
    plt.show()
```

```
numerical_df.skew(axis=0, numeric_only=True)
```

```
df = pd.concat([numerical_df, encoded_df], axis=1)  
df
```

✓ Question 6: Outlier Detection

Implement a method to detect outliers in the 'capital-gain' and 'capital-loss' columns.

a) What method did you choose and why? b) How many outliers did you detect? c) Propose a strategy for handling these outliers.

(a)**Method:** Z-score method using `scipy.stats.zscore()`

Reason:

- The Z-score allows us to quantify how extreme a data point is relative to the mean and standard deviation of the dataset.
- By setting a threshold ($|Z| > 3$), we can identify points that are significantly different from the rest of the data, which are considered outliers.

```
import scipy
a = df["capital-gain"][np.abs(scipy.stats.zscore(df["capital-gain"])) > 3]
a.hist()
```

```
b = df["capital-loss"][np.abs(scipy.stats.zscore(df["capital-loss"])) > 3]
b.hist()
```



```
print(len(a), len(b))
```

```
2100 1518
```

```
sns.boxplot(data=df[["capital-gain", "capital-loss"]])  
plt.show()
```

```

# TODO: Implement outlier detection for 'capital-gain' and 'capital-loss'
# Hint: Use the zscore() function from the scipy.stats module to detect outlier

# TODO: Visualize outliers
# Hint: A certain kind of plot is usually used to visualise outliers. Use the s
def handle_outliers(df, columns, quantile=0.95):
    """
    Handle outliers using winsorization.
    """
    for column in columns:
        lower_bound = df[column].quantile(1 - quantile)
        upper_bound = df[column].quantile(quantile)

        # Apply winsorization
        df[column] = np.where(df[column] < lower_bound, lower_bound, df[column])
        df[column] = np.where(df[column] > upper_bound, upper_bound, df[column])

    return df

# Call the function to handle outliers in 'capital-gain' and 'capital-loss'
df = handle_outliers(df, ['capital-gain', 'capital-loss'])

sns.boxplot(data=df[["capital-gain", "capital-loss"]])
plt.show()

```

✓ Question 7: Correlation Analysis

Perform a correlation analysis on the numerical variables.

a) Create a heatmap of the correlation matrix.

b) Identify any highly correlated pairs of features. How might this impact a machine learning model?

```
# TODO: Compute correlation matrix
# Hint: Use the corr() method on the DataFrame
corr = df.corr()
sns.heatmap(df[numerical].corr(), annot=True)
plt.show()
# TODO: Create a heatmap
# Hint: Use the heatmap() function from the seaborn library
high_corr = np.where((np.abs(corr) > 0.5) & (np.abs(corr) < 1))
l = []
for i in zip(*high_corr):
    l.append([df.columns[i[0]], df.columns[i[1]]])
pd.DataFrame(l, columns=["col1", "col2"])

# TODO: Identify and discuss highly correlated pairs
# Hint: You can use the where() and stack() methods to identify the highly corr
```

1. **Age and Marital Status (Never-married):** Younger people are more likely to be never married.
2. **Race (Black) and Race (White):** These race indicators are mutually exclusive.
3. **Marital Status (Married-civ-spouse) and Relationship (Not-in-family):** Married individuals are less likely to be classified as "Not-in-family."
4. **Marital Status (Married-civ-spouse) and Marital Status (Never-married):** These categories are mutually exclusive.

✓ Effects on the ML Model:

- **Overfitting:** Highly correlated variables might lead to overfitting if the model becomes too reliant on redundant information. This can reduce the model's performance on unseen data.
- **Feature Selection:** Feature selection techniques (e.g., mutual information or correlation thresholding) can help reduce the number of features and remove redundant information, making the model more interpretable and generalizable.
- **Multicollinearity:** Correlated variables can cause multicollinearity, making it harder for models (especially linear models) to identify the true effect of each feature. Regularization or feature reduction techniques can mitigate this.

#

Double-click (or enter) to edit

✓ Question 8: Class Imbalance

Investigate whether there is a class imbalance in the target variable ('income').

a) Calculate the proportion of each class in the target variable.

```
# TODO: Calculate and display class proportions
df["income"].value_counts(normalize=True)
```

✓ Question 9: Data Scaling

Implement feature scaling on the numerical variables.

a) Choose a scaling method (e.g., StandardScaler, MinMaxScaler) and explain your choice.

b) Apply the scaling to the numerical features.

```
# TODO: Choose and implement a scaling method
# Hint: The Scaler classes can be found in sklearn.preprocessing module
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
df[numerical] = scaler.fit_transform(df[numerical])
df
```

✓ Question 10: Exploratory Data Analysis

Perform exploratory data analysis to gain insights into the relationship between features and the target variable.

a) Create at least three different types of plots that reveal interesting patterns or relationships in the data. For instance, scatter plot, box plot, histogram and so on.

b) Explain your findings from each plot.

```
# TODO: Create at least three informative plots
```

```
# Example:
```

```
plt.figure(figsize=(12, 6))
```

```
sns.scatterplot(data=df, x='age', y='hours-per-week', hue='income')
```

```
plt.title('Scatter Plot: Age vs. Hours per Week by Income')
```

```
plt.show()
```

```
# Create 2 more such plots of your own
```

```
# TODO: Explain insights gained from each plot
```

- The scatter plot shows the relationship between age and hours worked per week, colored by income level.
- There appears to be a clustering of points where individuals earning more than \$50K tend to work more hours per week, especially in older age groups.
- Younger individuals tend to work fewer hours on average, regardless of income level.

```
#histogram for Age
plt.figure(figsize=(12, 6))
sns.histplot(df['age'], bins=20, kde=True)
plt.title('Histogram: Distribution of Age')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```

- The histogram shows the distribution of ages in the dataset.
- The age distribution appears to be roughly normal with a slight right skew, indicating that there are more younger individuals compared to older ones.
- There is a noticeable peak around ages 30-40, suggesting that this age group is well-represented in the dataset.

```
#bar plot for Average Hours per Week by Income
plt.figure(figsize=(12, 6))
average_hours = df.groupby('income')['hours-per-week'].mean().reset_index()
sns.barplot(data=average_hours, x='income', y='hours-per-week', palette='pastel')
plt.title('Bar Plot: Average Hours per Week by Income')
plt.xlabel('Income')
plt.ylabel('Average Hours per Week')
plt.show()
```


- The bar plot shows the average number of hours worked per week for each income class.
- Individuals earning more than 50K tend to work more hours on average compared to those earning less than or equal to 50K.
- This suggests that higher income may be associated with higher work commitments, which could be a factor in income classification.

```
#pair plot for selected features colored by Income
selected_features = df[['age', 'hours-per-week', 'capital-gain', 'income']]
sns.pairplot(selected_features, hue='income', palette='coolwarm')
plt.show()
```

- Age vs. Capital Gain: Older individuals tend to have higher capital gains, especially in the '>50K' income group.
- Hours per Week vs. Capital Gain: There is a noticeable trend indicating that individuals who work more hours also tend to have higher capital gains.
- Age vs. Hours per Week: The relationship shows that older individuals generally work more hours, which aligns with previous findings.